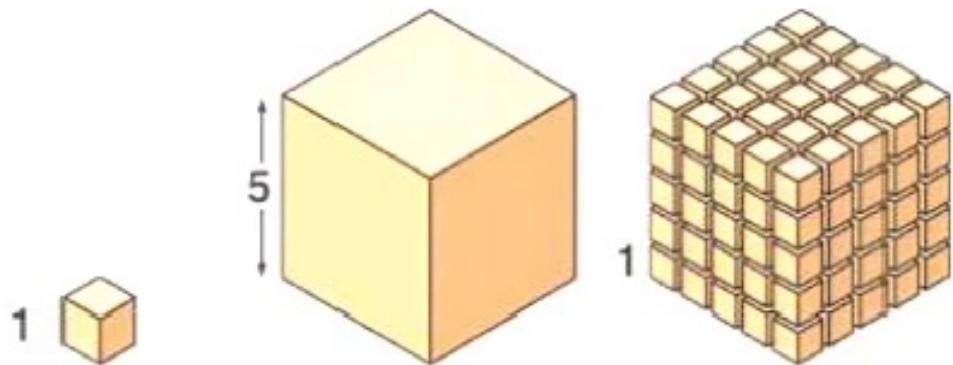
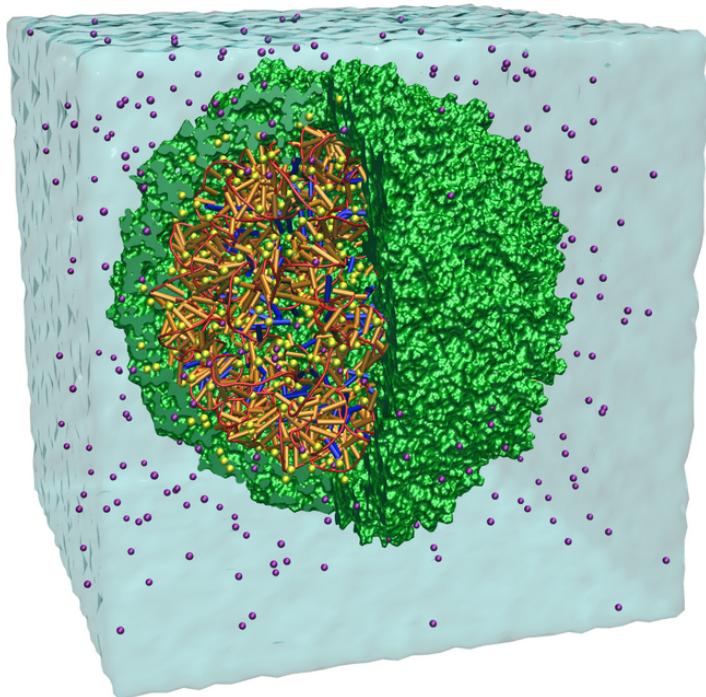


# High Performance Computing in Web Browsers

CE Seminar WT14/15

Henning Lohse

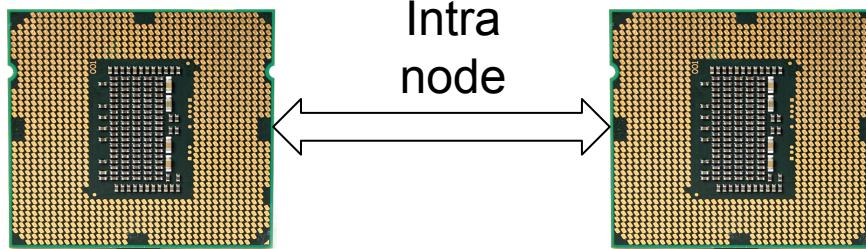
# High Performance Computing



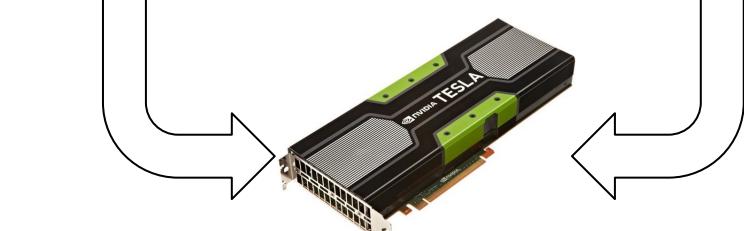
voltagegate.scientopia.org

# High Performance Computing

Processor

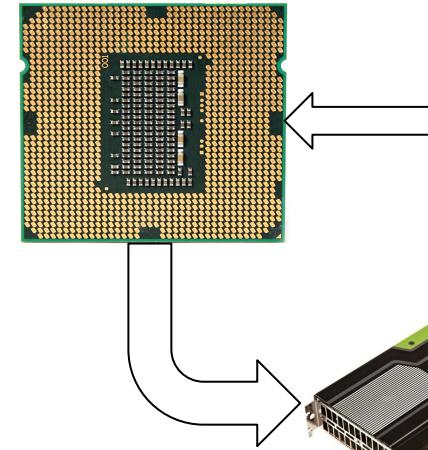


Intra  
node



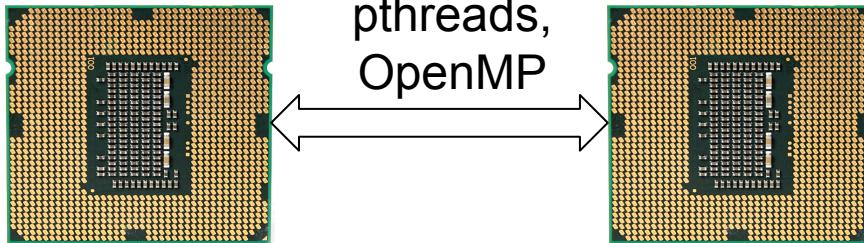
Processor  
GPU  
(coproc)

Inter  
node



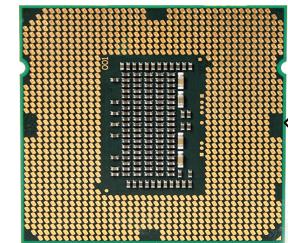
# High Performance Computing

C/C++

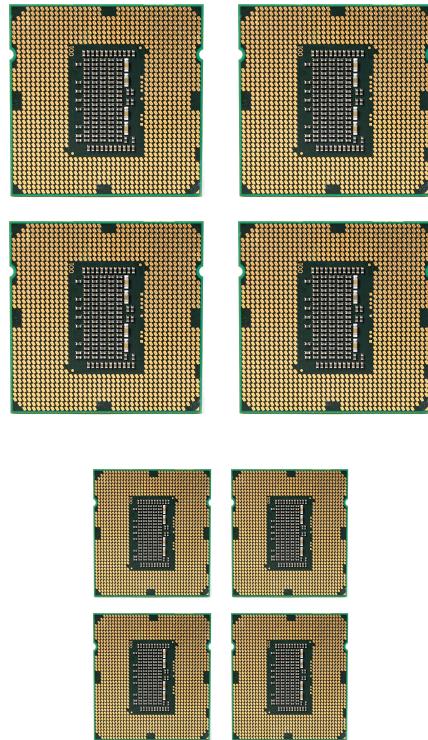


CUDA

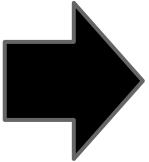
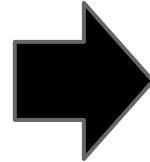
MPI,  
sockets



# Consumer Electronics

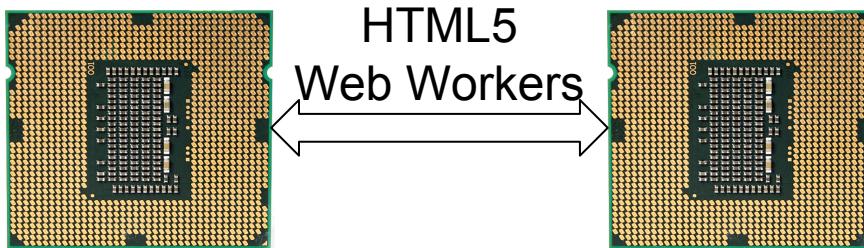


# Web Browsers



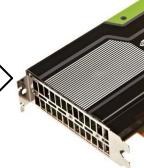
# HPC in Web Browsers

JavaScript, asm.js



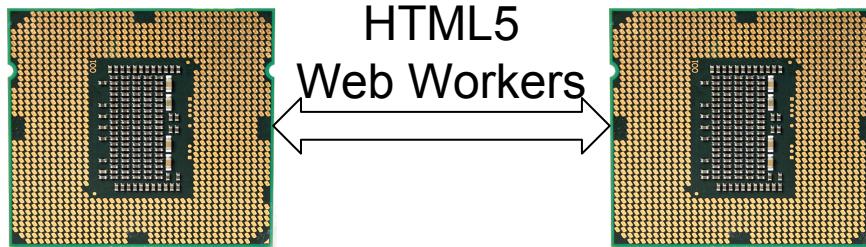
WebCL,  
WebGL

WebRTC  
Data-  
Channel



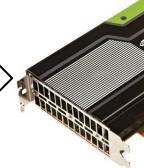
# JavaScript

JavaScript, asm.js

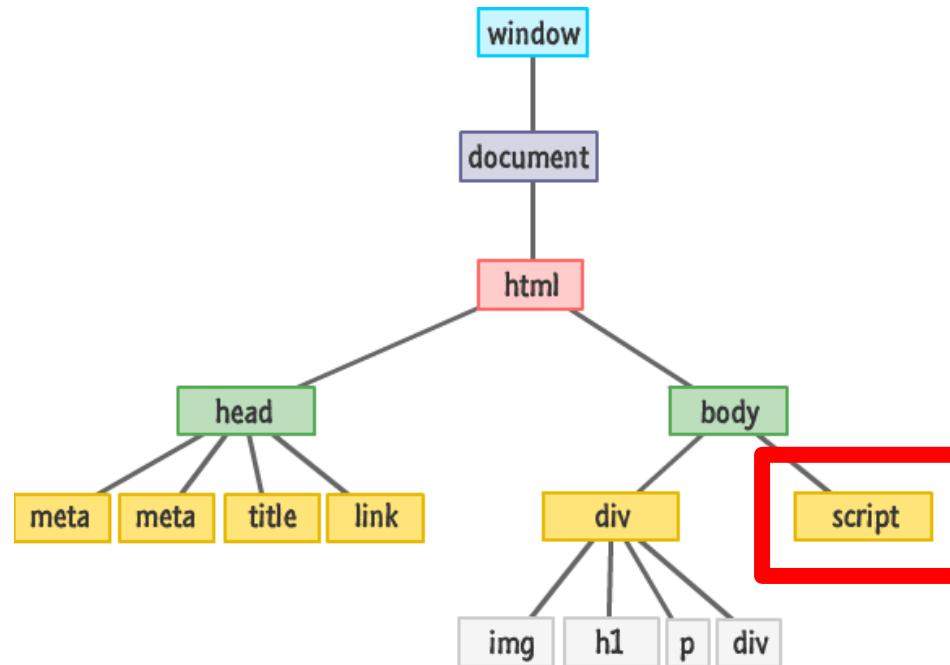


WebCL,  
WebGL

WebRTC  
Data-  
Channel



# JavaScript in HTML DOM



# JavaScript Object Notation: JSON

## JavaScript

- Dynamically typed
- Object-oriented
- Classless
- Prototypes

```
var prototype_object = {  
    "Herausgeber":      "Xema",  
    "Nummer":           "1234-5678-9012-3456",  
    "Deckung":          2e+6,  
    "Inhaber": {  
        "Name":       "Max",  
        "maennlich": true,  
        "Hobbys":     ["Reiten", "Golfen"],  
        "Alter":      42,  
    }  
};
```

```
var copy_object = prototype_object;  
copy_object["Waehrung"] = "EURO";  
copy_object["Inhaber"]["Hobbys"].push("Lesen");
```

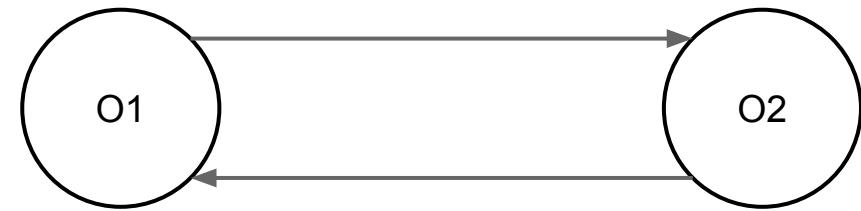
# Garbage Collection

Provided by Browsers

Reference Counting

Frees unused objects

- Implicit
- No leaks
- Undefined time



→ Mark and Sweep

# Performance Example

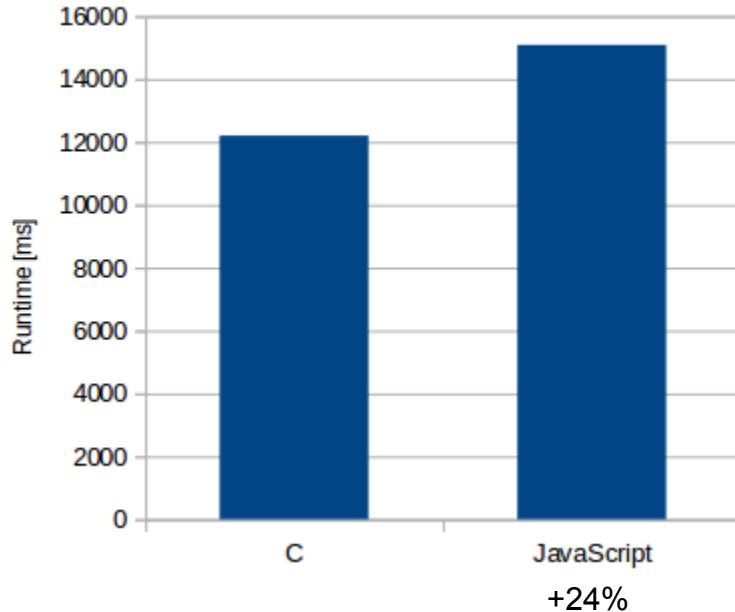
```
var primes = [];

for (var p = 2; p <= max; p++) {
    var is_prime = true;

    for (var i = 2; i <= Math.sqrt(max); i++)
        if (p % i == 0 && p != i) {
            is_prime = false;
            break;
        }

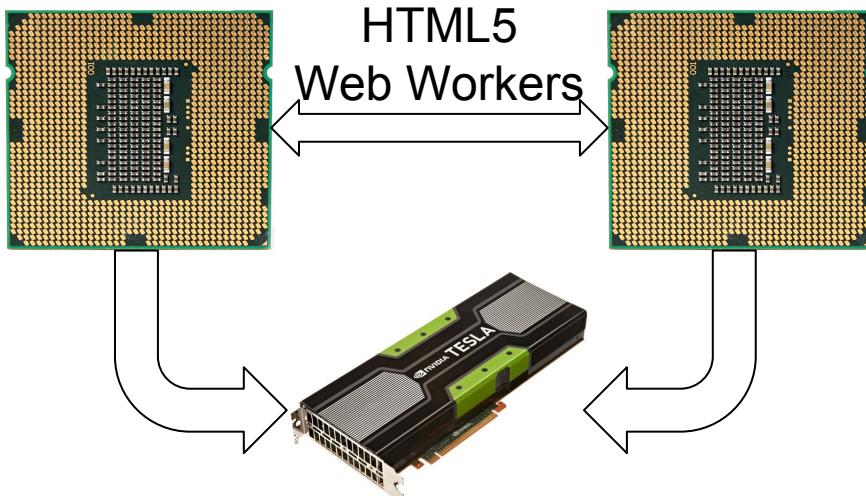
    primes[p] = is_prime;
}
```

Runtimes Finding First 10,000,000 Prime Numbers

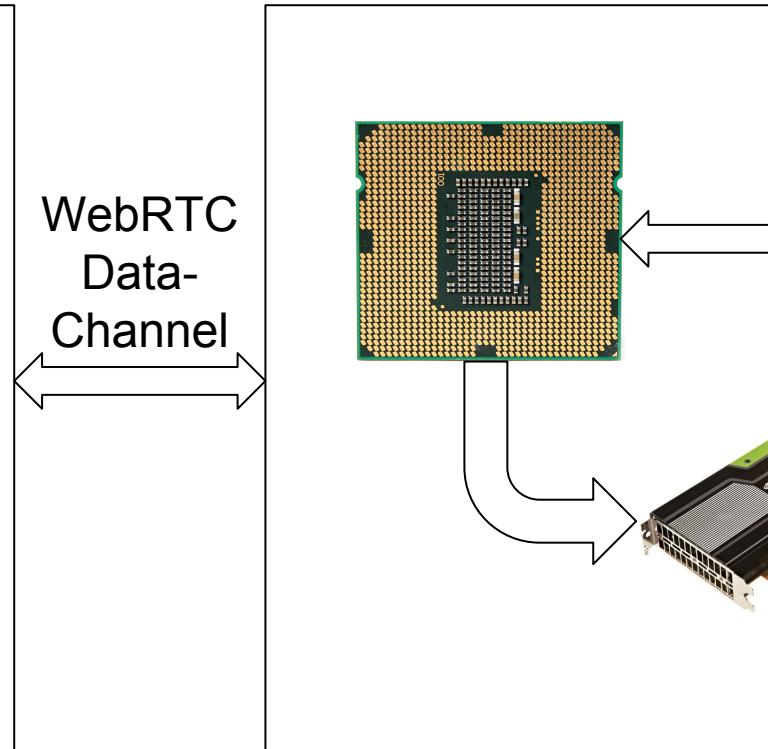


# asm.js

JavaScript, **asm.js**



WebRTC  
Data-  
Channel



# asm.js

“optimizable, low-level subset of JS” - Mozilla

```
var primes = [];

for (var p = 2; p <= max; p++) {
    var is_prime = true;

    for (var i = 2; i <= Math.sqrt(max); i++)
        if (p % i == 0 && p != i) {
            is_prime = false;
            break;
        }

    primes[p] = is_prime;
}
```

```
var primes = new Int32Array(max);

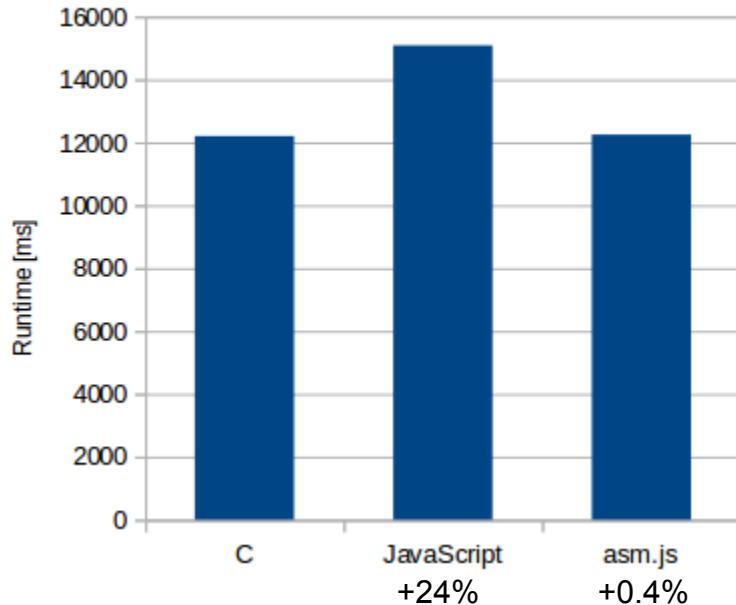
for (var p = (2|0); p <= max; p++) {
    var is_prime = (1|0);

    for (var i = (2|0); i <= (Math.sqrt(max)|0); i++)
        if (p % i == (0|0) && p != i) {
            is_prime = (0|0);
            break;
        }

    primes[p] = is_prime;
}
```

# asm.js

Runtimes Finding First 10,000,000 Prime Numbers



```
var primes = new Int32Array(max);

for (var p = (2|0); p <= max; p++) {
    var is_prime = (1|0);

    for (var i = (2|0); i <= (Math.sqrt(max)|0); i++)
        if (p % i == (0|0) && p != i) {
            is_prime = (0|0);
            break;
        }

    primes[p] = is_prime;
}
```

# LLVM



ISA of a virtual machine + compilation passes

C, C++, Java, Python, ... → LLVM IR → Binary

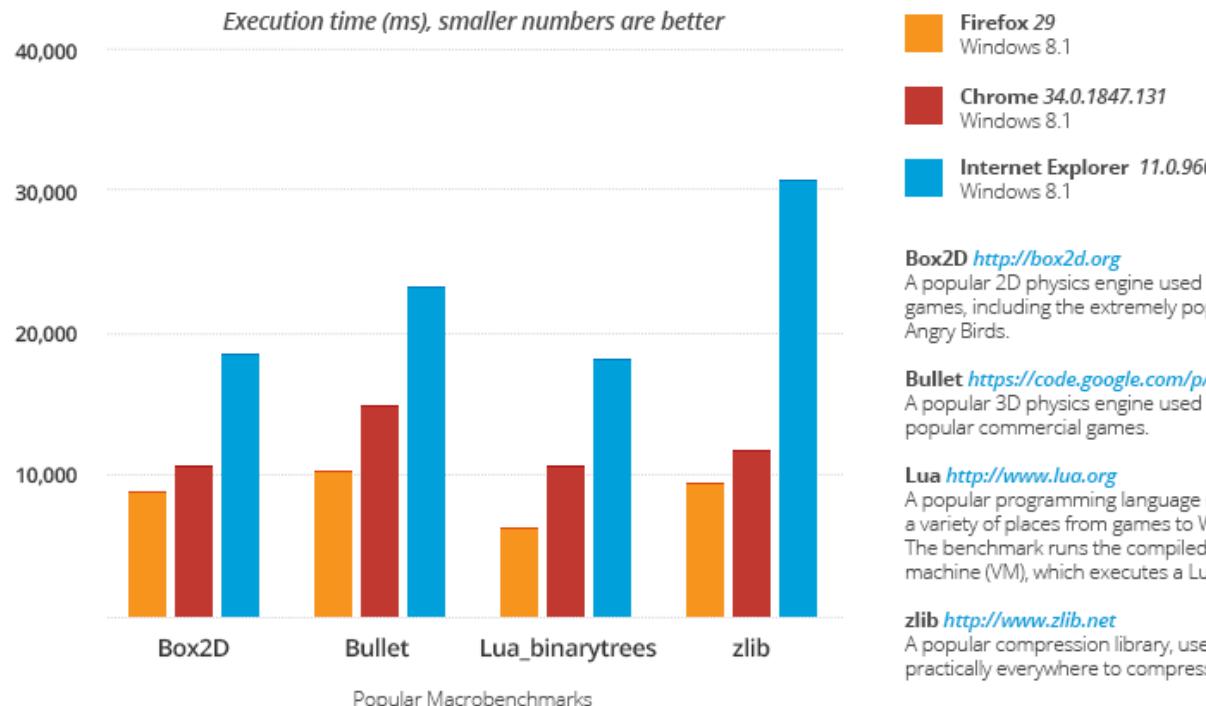
# Emscripten



C/C++ → LLVM IR → asm.js  
clang                    Emscripten

# asm.js Support

Benchmarks Showing asm.js Performance in Firefox vs Other Browsers (Windows)



# asm.js Summary

- JavaScript subset
- Annotation-based optimization detection
- Aims at near-native performance
- Compilable from C/C++
- Growing browser support

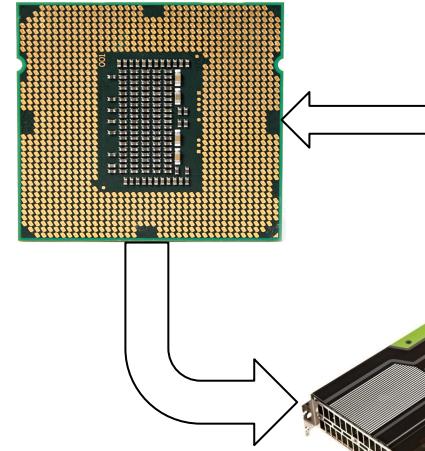
# HTML5 Web Workers

JavaScript, asm.js



WebCL,  
WebGL

WebRTC  
Data-  
Channel



# HTML5 Web Workers

Threads communicating via Message Passing

<script>

```
var worker = new Worker("worker_script.js");
```

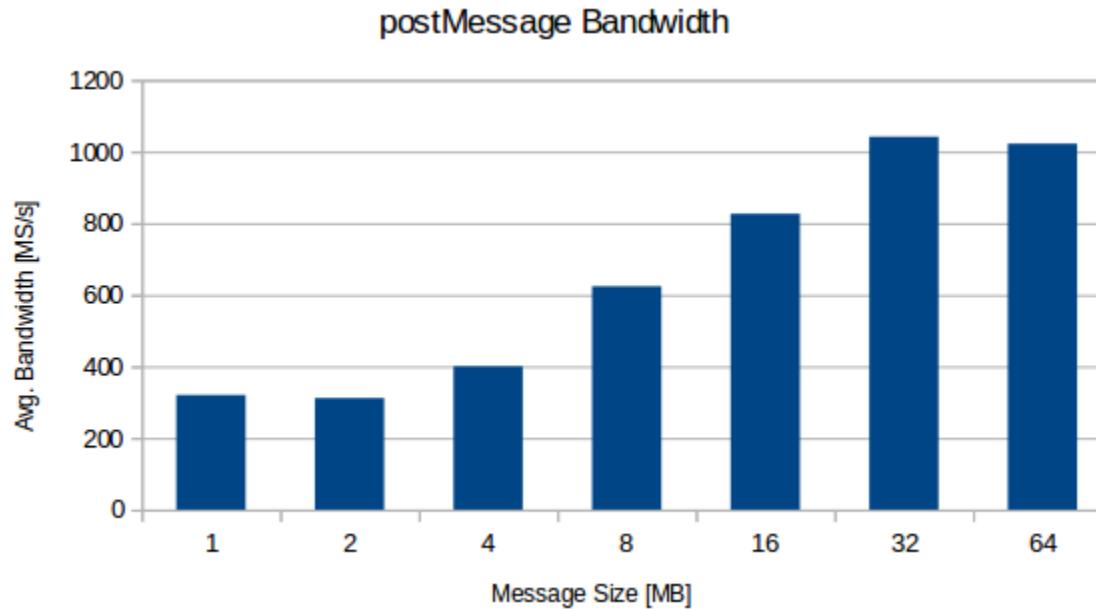
```
worker.addEventListener("message", function(e) {  
    console.log(e.data);  
}, false);
```

```
worker.postMessage("Hello!");  
</script>
```

worker\_script.js:

```
self.addEventListener("message", function(e) {  
    // Async computations go here  
    self.postMessage(e.data);  
}, false);
```

# postMessage: Structured Cloning



Garbage collection!

# postMessage: Transferable Objects

```
var array = new ArrayBuffer(1024); // 1kB  
worker.postMessage(array.buffer, [array.buffer]);
```

Latency: 53µs

Transferred data switch contexts!

# HTML5 Web Workers Summary

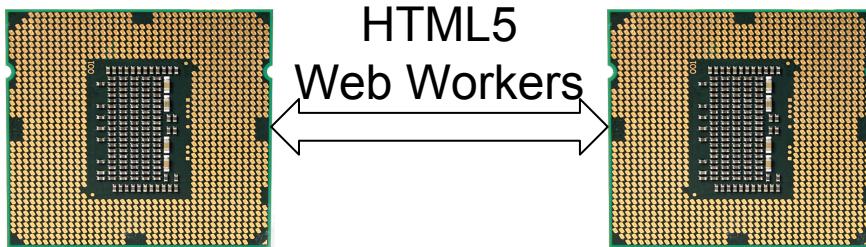
- No Shared Memory; Message Passing
- Cloning: Beware of GC, JSON
- Transfer: Beware of data context, JSON

Transferable objects support:



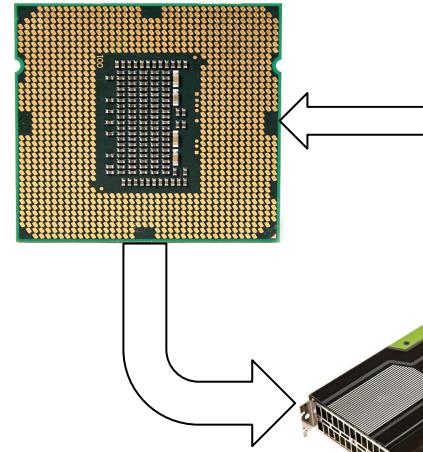
# WebRTC DataChannel

JavaScript, asm.js



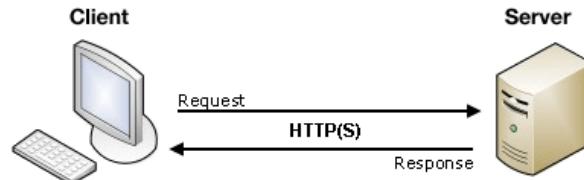
WebCL,  
WebGL

**WebRTC  
Data-  
Channel**



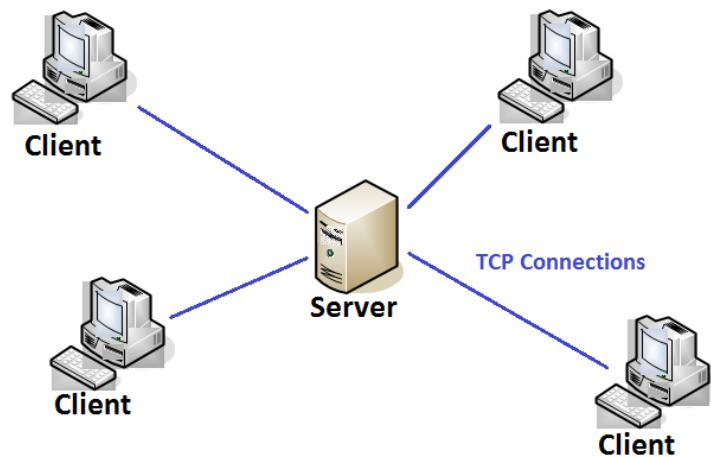
# HTTP

- Stateless
- GET, POST, ...
- Cookies

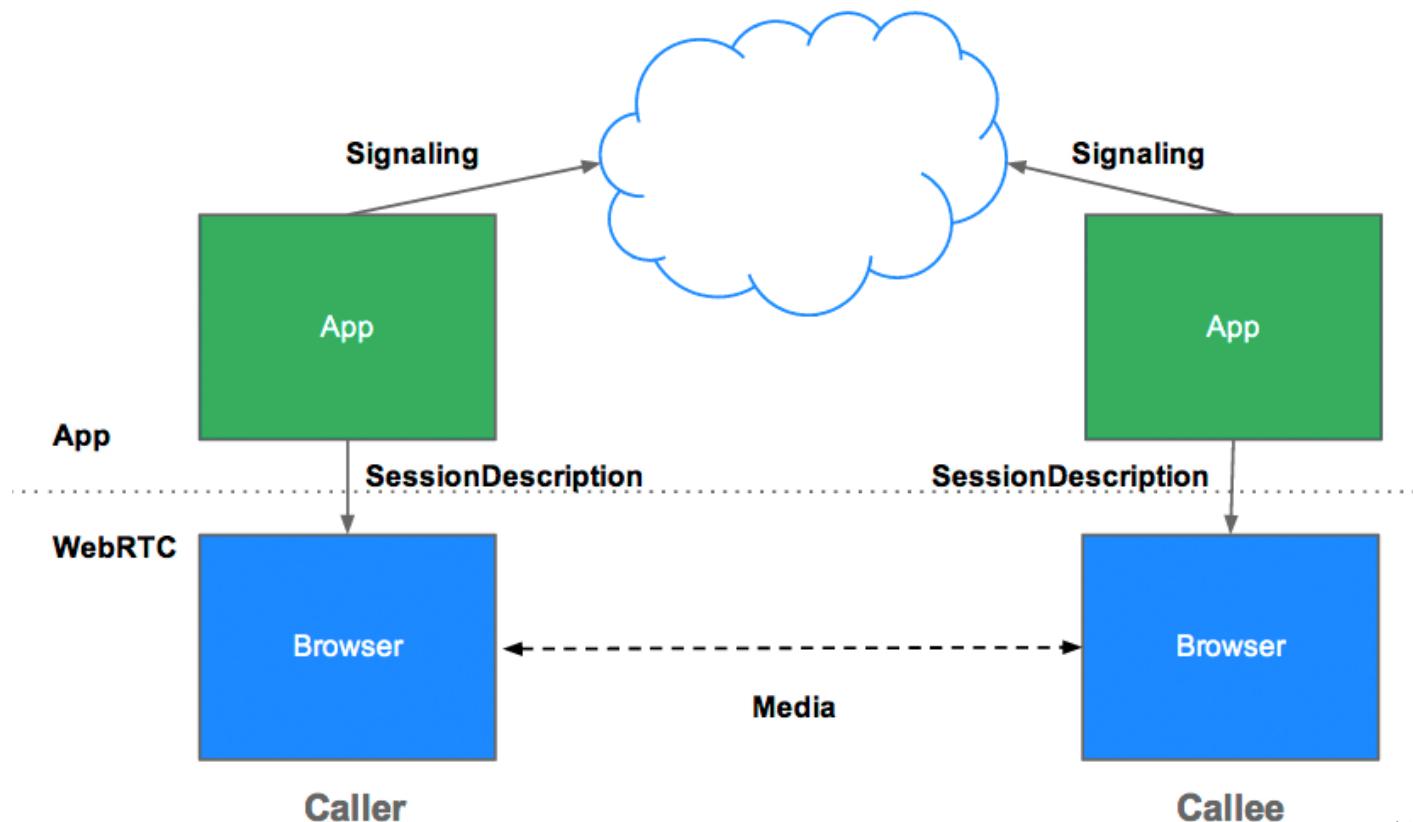


# HTML5 WebSocket

- Client server arch.
- TCP only

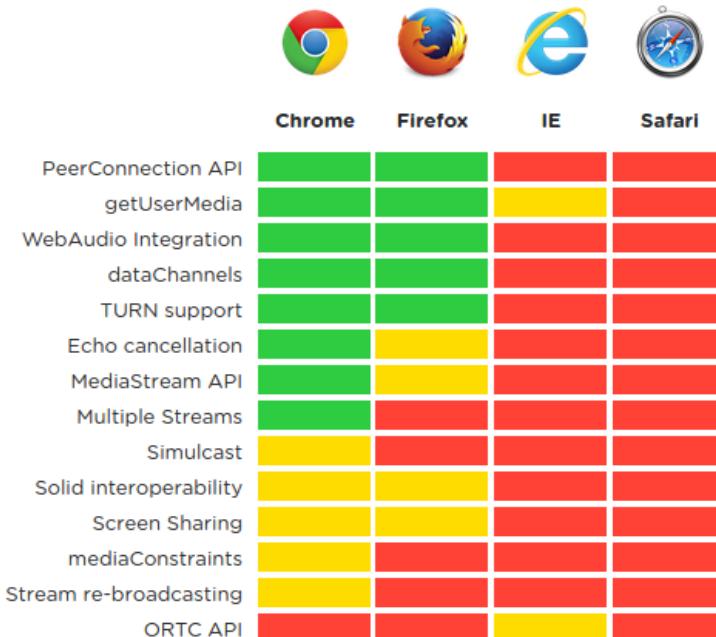


# WebRTC



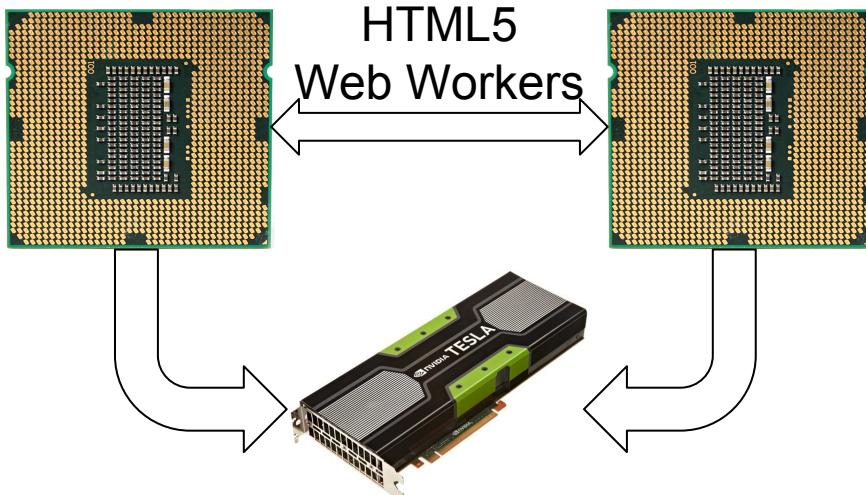
# WebRTC DataChannel

|                           | TCP           | UDP              | SCTP             |
|---------------------------|---------------|------------------|------------------|
| <b>Reliability</b>        | reliable      | unreliable       | configurable     |
| <b>Delivery</b>           | ordered       | unordered        | configurable     |
| <b>Transmission</b>       | byte-oriented | message-oriented | message-oriented |
| <b>Flow control</b>       | yes           | no               | yes              |
| <b>Congestion control</b> | yes           | no               | yes              |

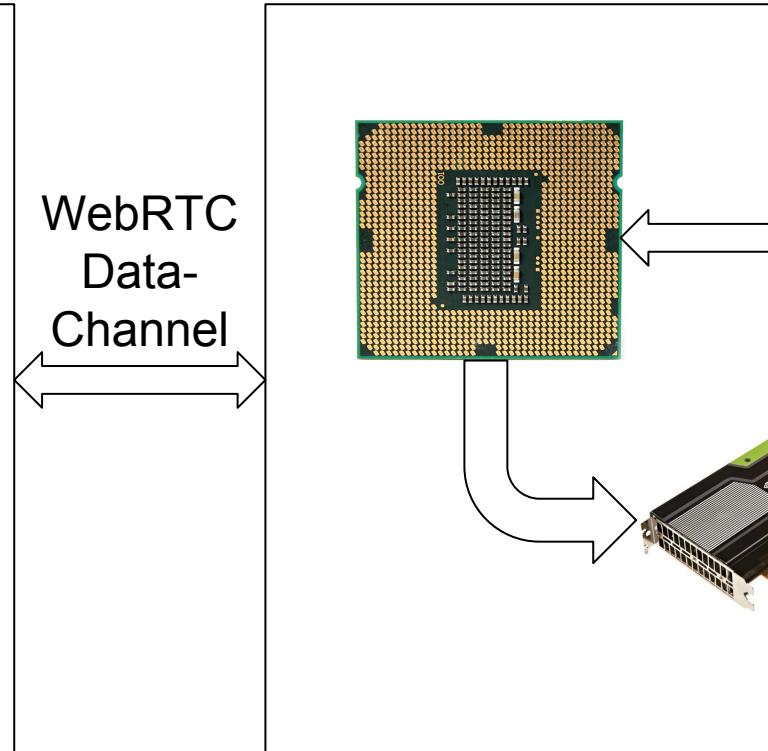


# WebCL, WebGL

JavaScript, asm.js

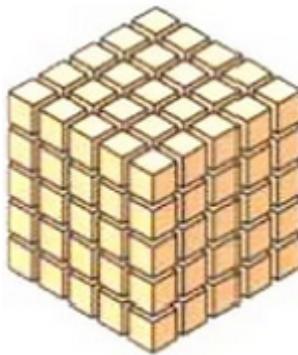
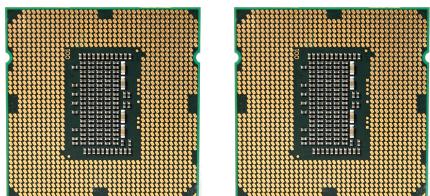


WebRTC  
Data-  
Channel



# GPU Computing

OpenCL/WebCL

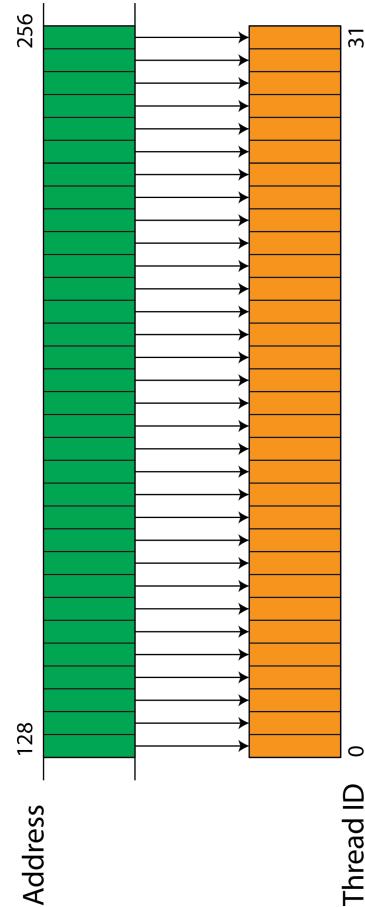


CUDA,  
OpenCL/WebCL,  
OpenGL/WebGL

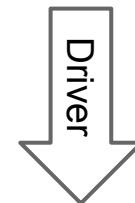


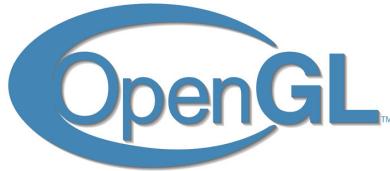


- Like OpenCL
  - Hardware exposure
  - IEEE 754 float
  - Heterogeneous
- 
- Drivers
  - Adaption



```
“__kernel void vectorAdd
    __global const float* x,
    __global const float* y,
    __global float* restrict z)
{
    int index = get_global_id(0);
    z[index] = x[index] + y[index];
}”
```





# Compute Shaders

- Since 4.3 (ES 3.1)
- GLSL, adaption
- Graphics abstraction
- No IEEE 754 float
- WebGL 1.0/2.0

```
layout (local_size_x = 16, local_size_y = 16) in;  
uniform readonly image2D fromTex;  
uniform writeonly image2D toTex;  
  
void main() {  
    ivec2 texelCoords = ivec2(gl_GlobalInvocationID.xy);  
    vec4 pixel = imageLoad(fromTex, texelCoords);  
    pixel.rg = pixel.gr;  
    imageStore(toTex, texelCoords, pixel);  
}
```



fullscreen Pause Resume Quit

# Summary

asm.js

JS subset, near-native speed

Web Workers

Messages: Cloning vs. transfer

DataChannel

Signalling, TCP/UDP/SCTP

WebCL/GL CS

Specs vs. adaption

# Summary



asm.js



Web Workers



DataChannel



WebCL/GL CS

